

Clustering

- **Review of K-means algorithm**

Algorithm: *k*-means. The *k*-means algorithm for partitioning, where each cluster's center is represented by the mean value of the objects in the cluster.

Input:

- *k*: the number of clusters,
- *D*: a data set containing *n* objects.

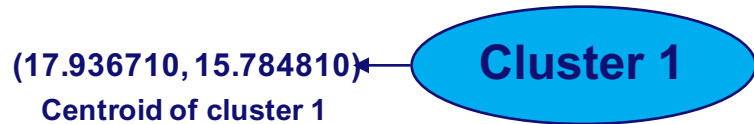
Output: A set of *k* clusters.

Method:

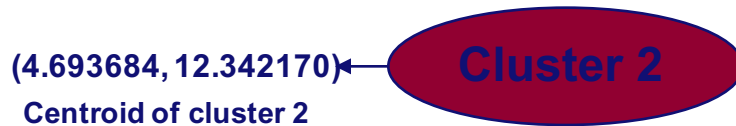
- (1) arbitrarily choose *k* objects from *D* as the initial cluster centers;
- (2) **repeat**
- (3) (re)assign each object to the cluster to which the object is the most similar,
 based on the mean value of the objects in the cluster;
- (4) update the cluster means, that is, calculate the mean value of the objects for
 each cluster;
- (5) **until** no change;

Clustering

- An example implementation of k-means algorithm



Initialization



$$d(i, j) = \sqrt{(x_{i1} - x_{j1})^2 + (x_{i2} - x_{j2})^2 + \dots + (x_{ip} - x_{jp})^2}$$

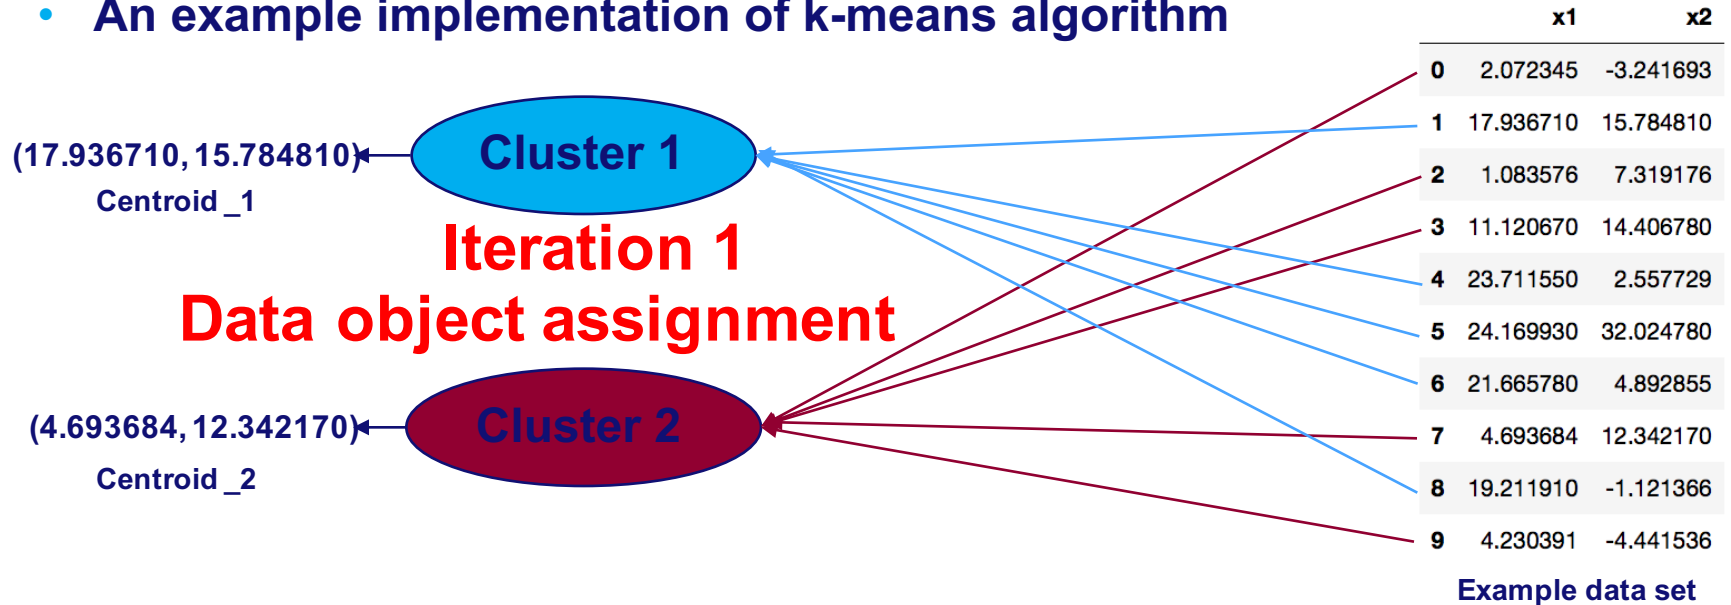
Euclidean distance

	x1	x2
0	2.072345	-3.241693
1	17.936710	15.784810
2	1.083576	7.319176
3	11.120670	14.406780
4	23.711550	2.557729
5	24.169930	32.024780
6	21.665780	4.892855
7	4.693684	12.342170
8	19.211910	-1.121366
9	4.230391	-4.441536

Example data set

Clustering

- An example implementation of k-means algorithm



Clustering

- An example implementation of k-means algorithm




Why the data item with index '9' is assigned to cluster 2 ???

$d(\text{object_9}, \text{centroid_1}) = ((4.230391 - 17.936710)^2 + (-4.441536 - 15.784810)^2)^{0.5} = 24.43293378$

$d(\text{object_9}, \text{centroid_2}) = ((4.230391 - 4.693684)^2 + (-4.441536 - 12.342170)^2)^{0.5} = 16.79009909$

$d(\text{object_9}, \text{centroid_2}) < d(\text{item_9}, \text{centroid_1}) \Rightarrow \text{item with index 9 is assigned to cluster 2}$



Example data set

$$d(i, j) = \sqrt{(x_{i1} - x_{j1})^2 + (x_{i2} - x_{j2})^2 + \dots + (x_{ip} - x_{jp})^2}$$

Euclidean distance

Clustering

- An example implementation of k-means algorithm

(17.936710, 15.784810)



(21.339176, 10.8277616)

Centroid_1

Cluster 1

Iteration 1

Update each centroid

(4.6401332, 5.2769794)



Centroid_2

Cluster 2

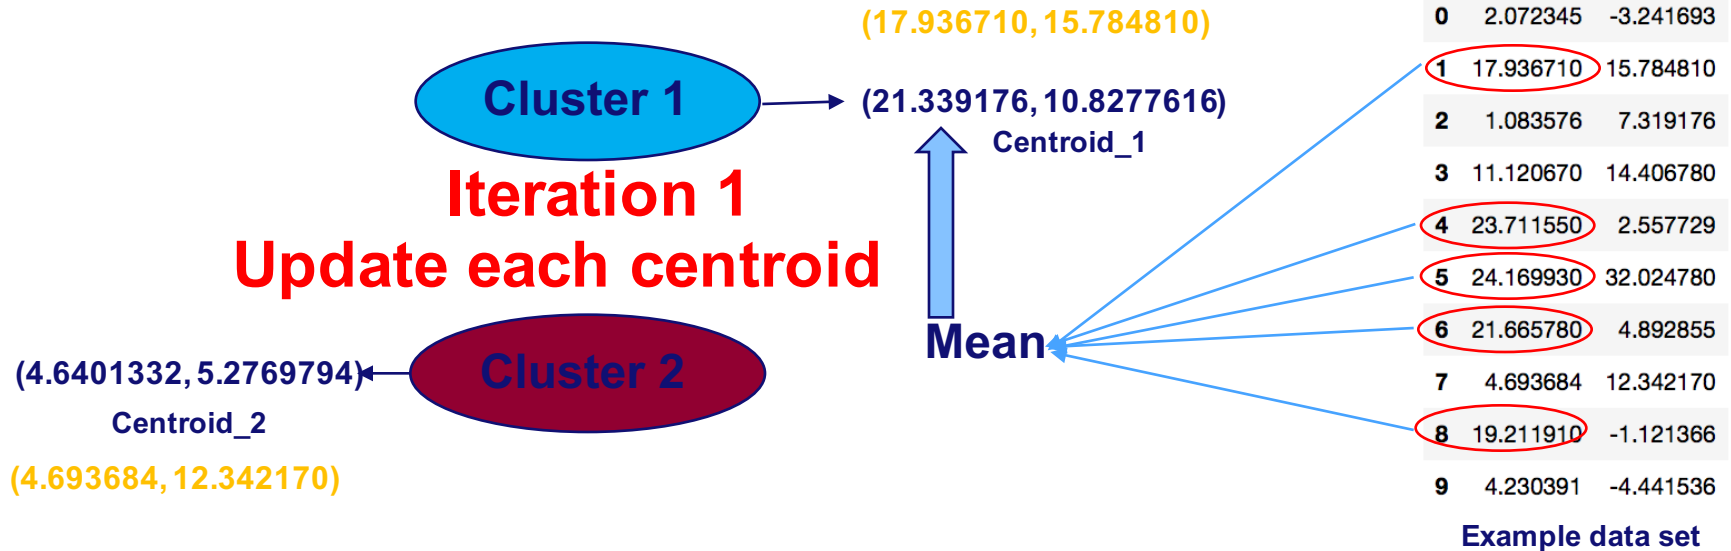
(4.693684, 12.342170)

	x1	x2
0	2.072345	-3.241693
1	17.936710	15.784810
2	1.083576	7.319176
3	11.120670	14.406780
4	23.711550	2.557729
5	24.169930	32.024780
6	21.665780	4.892855
7	4.693684	12.342170
8	19.211910	-1.121366
9	4.230391	-4.441536

Example data set

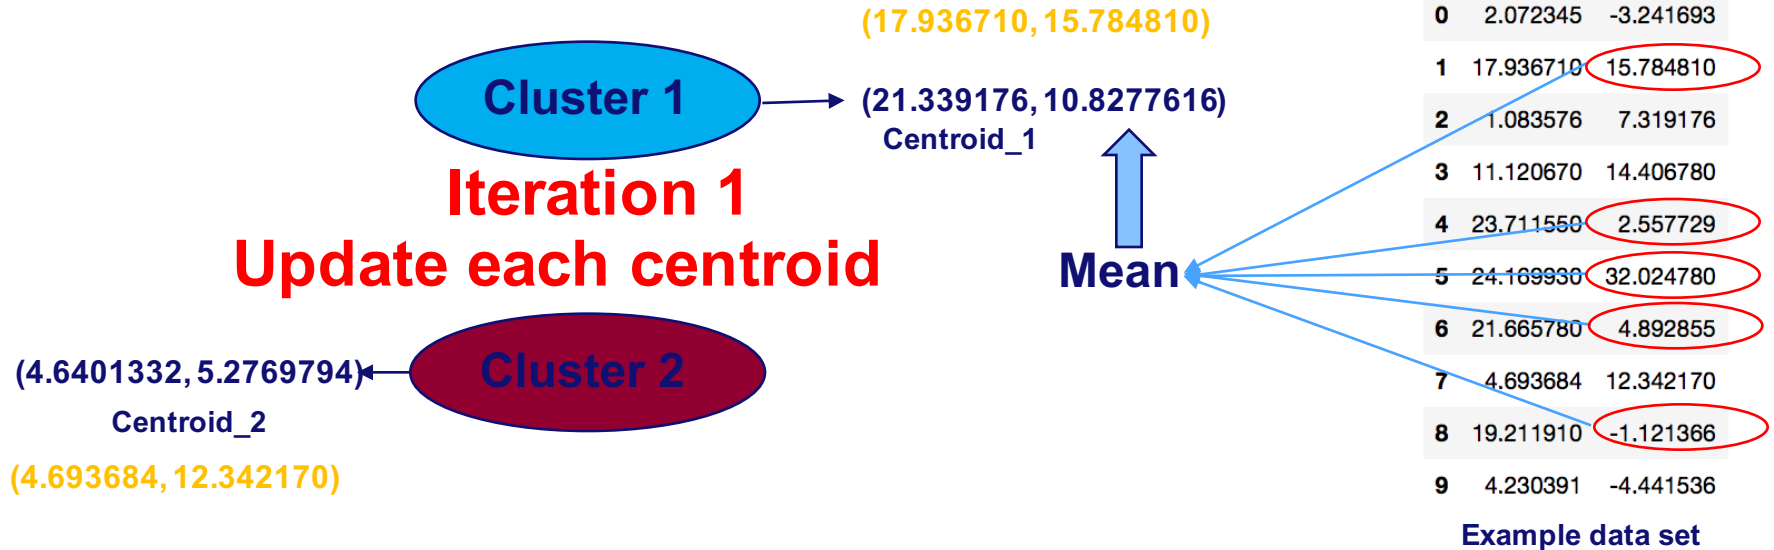
Clustering

- An example implementation of k-means algorithm



Clustering

- An example implementation of k-means algorithm



Clustering

- An example implementation of k-means algorithm

(17.936710, 15.784810)



(21.339176, 10.8277616)

Centroid_1

Cluster 1

Judge if centroids changed

(4.6401332, 5.2769794)



Centroid_2

Cluster 2

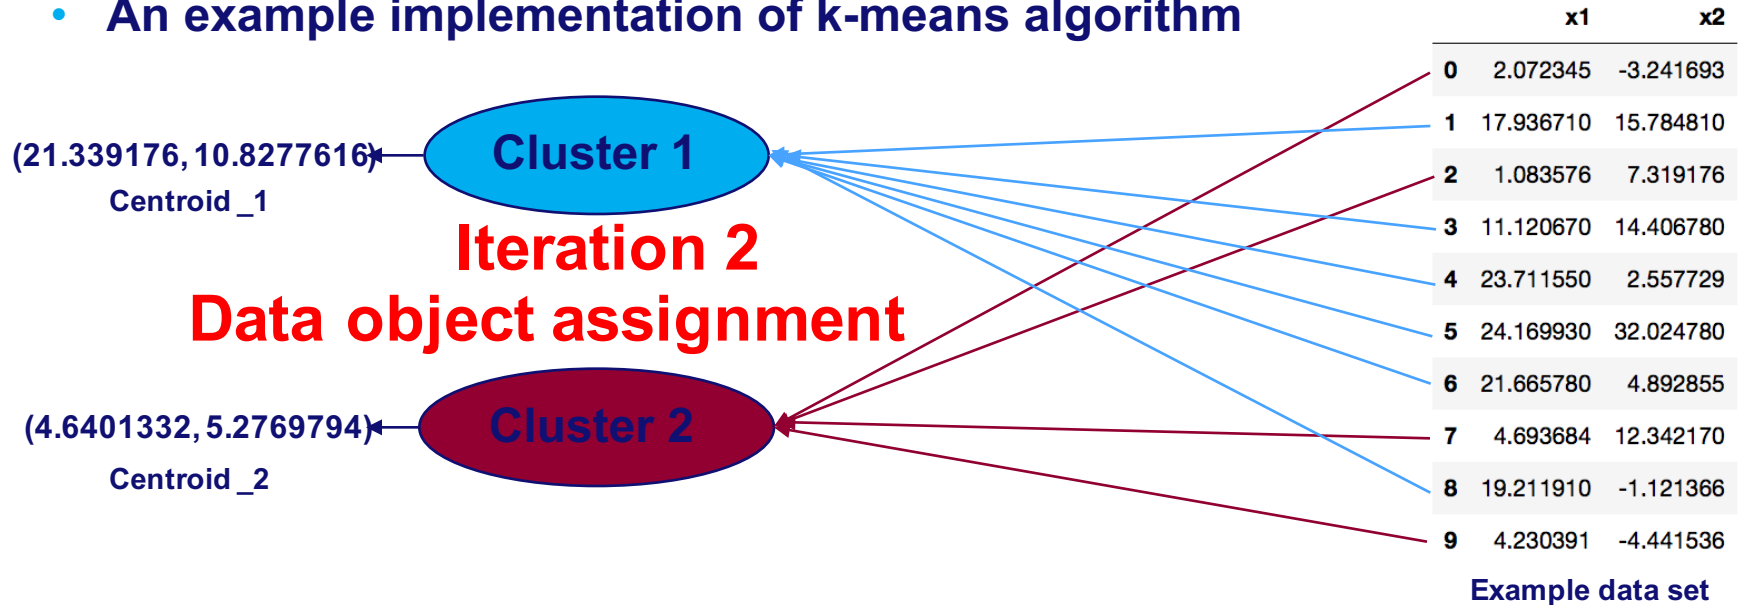
(4.693684, 12.342170)

	x1	x2
0	2.072345	-3.241693
1	17.936710	15.784810
2	1.083576	7.319176
3	11.120670	14.406780
4	23.711550	2.557729
5	24.169930	32.024780
6	21.665780	4.892855
7	4.693684	12.342170
8	19.211910	-1.121366
9	4.230391	-4.441536

Example data set

Clustering

- An example implementation of k-means algorithm



Clustering

- An example implementation of k-means algorithm

(21.339176, 10.8277616)

(19.636092, 11.4242647)

Centroid_1

Cluster 1

Iteration 2

Update each centroid

(4.6401332, 5.2769794)

(3.019999, 2.99452925)

Centroid_2

Cluster 2

	x1	x2
0	2.072345	-3.241693
1	17.936710	15.784810
2	1.083576	7.319176
3	11.120670	14.406780
4	23.711550	2.557729
5	24.169930	32.024780
6	21.665780	4.892855
7	4.693684	12.342170
8	19.211910	-1.121366
9	4.230391	-4.441536

Example data set

Clustering

- An example implementation of k-means algorithm

(21.339176, 10.8277616)

(19.636092, 11.4242647)

Centroid_1

Cluster 1

Judge if centroids changed

(4.6401332, 5.2769794)

(3.019999, 2.99452925)

Centroid_2

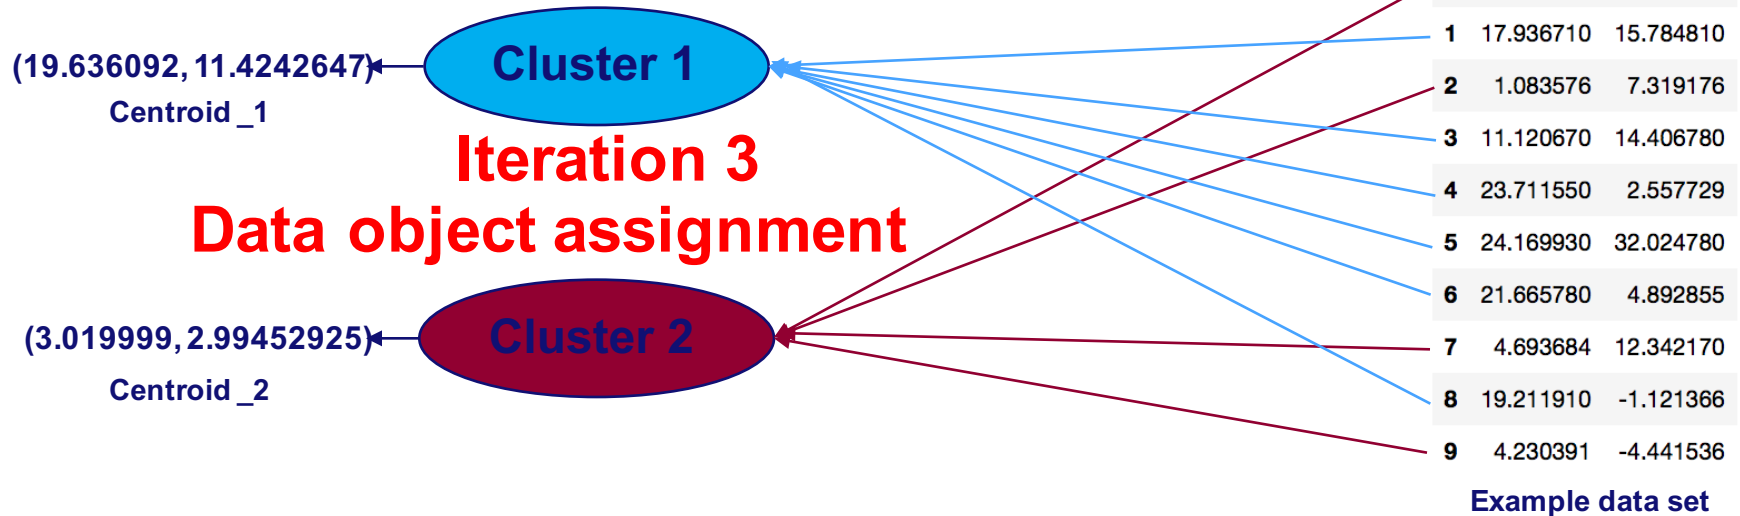
Cluster 2

	x1	x2
0	2.072345	-3.241693
1	17.936710	15.784810
2	1.083576	7.319176
3	11.120670	14.406780
4	23.711550	2.557729
5	24.169930	32.024780
6	21.665780	4.892855
7	4.693684	12.342170
8	19.211910	-1.121366
9	4.230391	-4.441536

Example data set

Clustering

- An example implementation of k-means algorithm



Clustering

- An example implementation of k-means algorithm

(19.636092, 11.4242647)

(19.636092, 11.4242647)

Centroid_1

Cluster 1

Iteration 3

Update each centroid

(3.019999, 2.99452925)

(3.019999, 2.99452925)

Centroid_2

Cluster 2

	x1	x2
0	2.072345	-3.241693
1	17.936710	15.784810
2	1.083576	7.319176
3	11.120670	14.406780
4	23.711550	2.557729
5	24.169930	32.024780
6	21.665780	4.892855
7	4.693684	12.342170
8	19.211910	-1.121366
9	4.230391	-4.441536

Example data set

Clustering

- An example implementation of k-means algorithm

(19.636092, 11.4242647)

(19.636092, 11.4242647)

Centroid_1

Cluster 1

Judge if centroids changed

(3.019999, 2.99452925)

(3.019999, 2.99452925)

Centroid_2

Cluster 2

	x1	x2
0	2.072345	-3.241693
1	17.936710	15.784810
2	1.083576	7.319176
3	11.120670	14.406780
4	23.711550	2.557729
5	24.169930	32.024780
6	21.665780	4.892855
7	4.693684	12.342170
8	19.211910	-1.121366
9	4.230391	-4.441536

Example data set

Clustering

- **Exercise 1: cluster data objects from data set 1 manually**
- **Conditions**
 - **K = 2** :- cluster 1 and cluster 2
 - Data object with index '4' as initial centroid for cluster 1
 - Data object with index '8' as initial centroid for cluster 2
 - Use Euclidean distance
- **You should calculate and output**
 - Centroids of the found clusters
 - Data objects for each of the two final clusters

	x1	x2
0	0.407503	15.297050
1	7.314846	3.309312
2	-3.438403	-12.025270
3	17.639350	-3.212345
4	4.415292	22.815550
5	11.941220	8.122487
6	0.725853	1.806819
7	8.185273	28.132600
8	-5.773587	1.024800
9	18.769430	24.169460

Data set 1

$$d(i, j) = \sqrt{(x_{i1} - x_{j1})^2 + (x_{i2} - x_{j2})^2 + \dots + (x_{ip} - x_{jp})^2}$$

Euclidean distance

Clustering

- **Some weaknesses of k-means algorithm**
 - Number of clusters needs to be decided beforehand
 - **Can only discover spherical clusters (compare to density-based methods)**
 - Sensitive to outliers (show this to you later)

Frequent Itemsets

TID	Set of items
0	bread, meat, wine
1	bread, meat
2	pizza, wine
3	bread, meat, pizza, wine

Set of transactions D

Minimum support_{count} = 2



Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets

Frequent Itemsets

- Basic ideas of Apriori algorithm
 - **Apriori rule:** all the non-empty **sub-itemsets** of frequent itemsets must **be frequent**.

Frequent Itemsets

TID	Set of items
0	bread, meat, wine
1	bread, meat
2	pizza, wine
3	bread, meat, pizza, wine

Set of transactions D

Minimum support_{count}
= 2



Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets

Itemset {bread, meat, wine}
is frequent



Sub-item set	Support	
{bread}	3	frequent
{meat}	3	frequent
{wine}	3	frequent
{bread, meat}	3	frequent
{bread, wine}	2	frequent
{meat, wine}	2	frequent




Frequent Itemsets

- **Basic ideas of Apriori algorithm**
 - **Apriori rule:** all the non-empty **sub-itemsets** of **frequent itemsets** should **be frequent**.
 - Use the set L_k of frequent itemsets with **length k** to search for both candidate set C_{k+1} of itemsets and the set L_{k+1} of frequent itemsets with length **$k+1$**

Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

L₁



First, Scan D to generate **L₁**


Frequent itemsets

Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets

First, Scan D to generate L_1



TID	Set of items
0	bread, meat, wine
1	bread, meat
2	pizza, wine
3	bread, meat, pizza, wine

Set of transactions D

min_sup represents
minimum support_{count}

compare support
to min_sup
Scan,
remove infrequent
ones

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3

L_1



Frequent Itemsets

- Exercise 2: manually generate L_1 from set S
 - Set minimum support_{count} to 2

Item set	Support
→ ?	

L_1


TID	Data items
1	A, B, E
2	C, A, D
3	C, B, D
4	C, A, B, E

Example data set S

Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	L_2 2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	L_3 2

Frequent itemsets




How to generate C_k from L_{k-1} , when $k \geq 2$?

Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets



How to generate C_k from L_{k-1} , when $k \geq 2$?

1. Keep items in itemsets in L_{k-1} in an ascending order according to their dictionary order

Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets

How to generate C_k from L_{k-1} , when $k \geq 2$?

1. Keep items in itemsets in L_{k-1} in an ascending order according to their dictionary order

Set of items	Support
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2

L_2

For generating C_3

bread < meat

bread < wine

meat < wine


pizza < wine



Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets



How to generate C_k from L_{k-1} , when $k \geq 2$?

2. Merge each possible pair of items from L_{k-1}

Let $(\{p_1, p_2, \dots, p_{k-2}, p_{k-1}\}, \{q_1, q_2, \dots, q_{k-2}, q_{k-1}\})$ be a pair of two itemsets from L_{k-1} .

-- if: $\{p_1, p_2, \dots, p_{k-2}\} = \{q_1, q_2, \dots, q_{k-2}\}$ and $p_{k-1} < q_{k-1}$

-- then merge them into: $\{p_1, p_2, \dots, p_{k-2}, p_{k-1}, q_{k-1}\}$


-- else: not merge them



Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets



How to generate C_k from L_{k-1} , where $k \geq 2$?

2. Merge each possible pair of items from L_{k-1}

Set of items	Support
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2

L_2



$(\{\text{bread, meat}\}, \{\text{bread, wine}\})$
 $\Rightarrow \{\text{bread, meat, wine}\}$

Because $\{\text{bread}\} = \{\text{bread}\}$ and
 $\text{meat} < \text{wine}$

Frequent Itemsets

Item set Support

{bread} 3

{meat} 3

{pizza} 2

{wine} 3

{bread, meat} 3

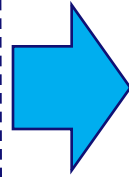
{bread, wine} 2

{meat, wine} 2

{pizza, wine} 2

{bread, meat, wine} 2

Frequent itemsets



How to generate C_k from L_{k-1} , where $k \geq 2$?

2. Merge each possible pair of items from L_{k-1}

Set of items Support

{bread, meat} 3

{bread, wine} 2

{meat, wine} 2

{pizza, wine} 2

L_2



$(\{\text{bread, wine}\}, \{\text{bread, meat}\})$
 $\nRightarrow \{\text{bread, meat, wine}\}$

Because **wine** > **meat**



Frequent Itemsets

Item set Support

{bread} 3

{meat} 3

{pizza} 2

{wine} 3

{bread, meat} 3

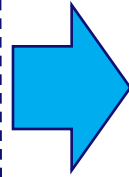
{bread, wine} 2

{meat, wine} 2

{pizza, wine} 2

{bread, meat, wine} 2

Frequent itemsets



How to generate C_k from L_{k-1} , where $k \geq 2$?

2. Merge each possible pair of items from L_{k-1}

Set of items Support

{bread, meat} 3

{bread, wine} 2

{meat, wine} 2

{pizza, wine} 2

L_2



$(\{\text{bread, wine}\}, \{\text{meat, wine}\})$
 $\nRightarrow \{\text{bread, meat, wine}\}$

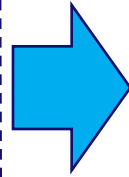
Because $\{\text{bread}\} \neq \{\text{meat}\}$



Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets



How to generate C_k from L_{k-1} , where $k \geq 2$?

2. Merge each possible pair of items from L_{k-1}

Set of items	Support
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2

L_2

merge itemsets
in L_2

Item set
{bread, meat, wine}

C_3

Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	L_2 2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	L_3 2

Frequent itemsets




Pre-pruning for C_k when $k \geq 2$ before scanning D

Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets




Pre-pruning for C_k when $k \geq 2$ before scanning D

Pruning rule: given an itemset $i_k \subseteq C_k$, if not all sub-itemsets of length $k-1$ of i_k are contained in L_{k-1} , then remove i_k from C_k

Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2

Frequent itemsets



Pre-pruning for C_k when $k \geq 2$ before scanning D

Set of items	Support
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2

L_2

The sub-itemsets of length 2 for {bread, meat, wine} include: {bread, meat}, {bread, wine}, {meat, wine}, which are all Included in L_2 . So keep {bread, meat, wine}.

Item set
{bread, meat, wine}


C_3



Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	L_2 2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	L_3 2

Frequent itemsets



Pre-pruning for C_k when $k \geq 2$ before scanning D





Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2


Frequent itemsets

Scan **D** to add supports for itemsets in **C_k**





Frequent Itemsets

Item set	Support
{bread}	3
{meat}	3
{pizza}	2
{wine}	3
{bread, meat}	3
{bread, wine}	2
{meat, wine}	2
{pizza, wine}	2
{bread, meat, wine}	2



Frequent itemsets

Remove infrequent itemsets and generate L_k



Frequent Itemsets

- Exercise 3: given L_2 for set S , show the three C_3 and L_3
 - Set minimum support_{count} to 2

TID	Data items
1	A, B, E
2	C, A, D
3	C, B, D
4	C, A, B, E

Data set S

Item set	Support
{A, B}	2
{A, C}	2
{A, E}	2
{B, C}	2
{B, E}	2
{C, D}	2

L_2

merge itemsets
in L_2

Item set	Support

C_3

pruning

Item set	Support

C_3

Scan S

Item set	Support

C_3

compare support
to min_sup

remove infrequent
ones

Item set	Support

L_3

Association Rules

- **T** is a set of transactions
- **I** is the set of all possible itemsets composed by items in **T**
- **A** \subseteq **I** and **B** \subseteq **I** are two itemsets/sub-itemsets from **T**
- **A** \Rightarrow **B** is an association rule

Association Rules

- Usually, we would like to discover the association rule $A \Rightarrow B$ of which the support and confidence are above certain levels.

Association Rules

- $support(A \Rightarrow B) = support(A \cup B) = \frac{support_{count}(A \cup B)}{|T|}$
- $confidence(A \Rightarrow B) = \frac{support(A \cup B)}{support(A)} = \frac{support_{count}(A \cup B)}{support_{count}(A)}$
- **min_sup** represents **minimum support** and **min_conf** represents **minimum confidence**
- **A => B** is a desired association rule if:
 - $support(A \Rightarrow B) \geq min_sup$ and $confidence(A \Rightarrow B) \geq min_conf$



Association Rules

- Set **min_sup** to 0.5, **min_conf** to 0.7, is **{bread} => {meat}** from **D** a desired association rule?

TID	Set of items
0	bread, meat, wine
1	bread, meat
2	pizza, wine
3	bread, meat, pizza, wine

Set of transactions D

Association Rules

- Set **min_sup** to 0.5, **min_conf** to 0.7, is **{bread} => {meat}** from **D** a desired association rule?

TID	Set of items
0	bread, meat, wine
1	bread, meat
2	pizza, wine
3	bread, meat, pizza, wine

Set of transactions **D**

$$\text{support}(\{\text{bread}\} \Rightarrow \{\text{meat}\}) = \frac{\text{support}_{\text{count}}(\{\text{bread, meat}\})}{|D|} = \frac{3}{4} = 0.75 > \text{min_sup}$$

$$\text{confidence}(\{\text{bread}\} \Rightarrow \{\text{meat}\}) = \frac{\text{support}_{\text{count}}(\{\text{bread, meat}\})}{\text{support}_{\text{count}}(\{\text{bread}\})} = \frac{3}{3} = 1 > \text{min_conf}$$



{bread} => {meat} is a desired association rule

Association Rules

- Usually, we use **lift** to evaluate the quality of the discovered association rule **$A \Rightarrow B$**

$$\text{lift}(A \Rightarrow B) = \frac{\text{support}(A \cup B)}{\text{support}(A) \cdot \text{support}(B)} = \frac{P(A \cup B)}{P(A) \cdot P(B)}$$

If $\text{lift}(A \Rightarrow B) \approx 1$ then A and B are independent

If $\text{lift}(A \Rightarrow B) \ll 1$ then A and B are negatively correlated

If $\text{lift}(A \Rightarrow B) \gg 1$ then A and B are positively correlated

Association Rules

- Evaluate the quality of the association rule $\{bread\} \Rightarrow \{meat\}$ by using **lift**

TID	Set of items
0	bread, meat, wine
1	bread, meat
2	pizza, wine
3	bread, meat, pizza, wine

Set of transactions D

$$lift(\{bread\} \Rightarrow \{meat\}) = \frac{support(\{bread, meat\})}{support(\{bread\}) \cdot support(\{meat\})} = \frac{(3/4)}{(3/4) \cdot (3/4)} = 1.33$$

Association Rules

- Exercise 3: judge if $\{A, B\} \Rightarrow \{E\}$, $\{A\} \Rightarrow \{B\}$ and $\{A\} \Rightarrow \{C\}$ are the desired association rules under minimum support 0.5 and minimum confidence 0.75? Also evaluate the quality of the desired rules.

TID	Data items
1	A, B, E
2	C, A, D
3	C, B, D
4	C, A, B, E

Example data set S