

Informed Search

Introduction to Artificial Intelligence

G. Lakemeyer

Winter Term 2018/19

Best-First Search

Search methods differ in their strategies which node to expand next

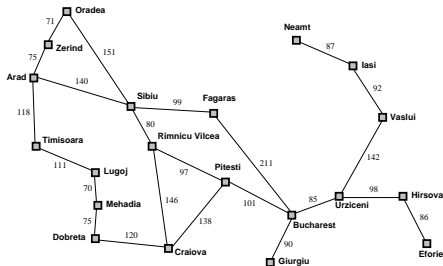
Uninformed search: fixed strategies without information about the cost from a given node to a goal.

Informed search: uses information about the cost from a given node to a goal in the form of an **evaluation function** f , assigning each node a real number.

Best-First Search: always expand the node with the “best” f -value.

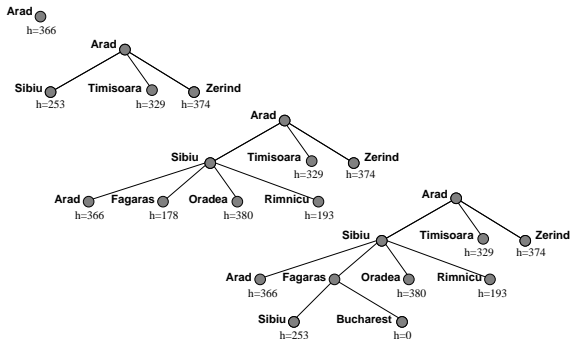
Greedy Search: $h(n)$ = estimated cost from state at node n to a goal state. Expand node n where $h(n)$ is minimal.
Use $f = h$.

Greedy Search Example



Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374



A*

combines uniform cost search with greedy search.

$g(n)$ = actual cost from the initial state to n .

$h(n)$ = estimated cost from n to the nearest goal.

$f(n) := g(n) + h(n)$.

$f(n)$ = is the estimated cost of the cheapest path which passes through n .

Let $h^*(n)$ be the actual cost of the optimal path from n to the nearest goal.

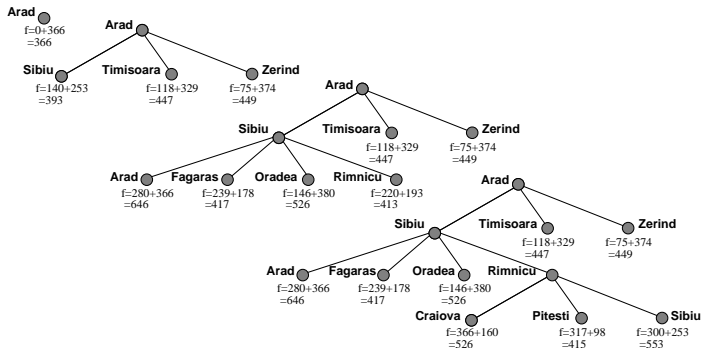
Admissible Heuristic

h is called **admissible** if we have for all n :

$$h(n) \leq h^*(n) .$$

We require for A* that h is admissible.

Example A*




Note: in the example f is **monotone** nondecreasing.
The following can always be guaranteed:

Path-Max Equation

Let n, n' be nodes, where is n parent of n' . Then let

$$f(n') = \max(f(n), g(n') + h(n')).$$

Contour Lines in A^*



Heuristic Function 1

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

h_1 = number of tiles in the wrong position.

h_2 = sum of the distances to the goal location for all tiles
(Manhattan Distance)

Heuristic Function 1

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5

Goal State

h_1 = number of tiles in the wrong position.

h_2 = sum of the distances to the goal location for all tiles
(Manhattan Distance)

Effective branching factor b^* : If A^* generates N nodes with solution depth d , then b^* is the branching factor of a uniform tree of depth d with $N + 1$ nodes, i.e.

$$N + 1 = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$$

b^* is a measure for the goodness of h : the closer b^* is to 1 the better.

Heuristic Function 2

d	Search Cost			Effective Branching Factor		
	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	364404	227	73	2.78	1.42	1.24
14	3473941	539	113	2.83	1.44	1.23
16	—	1301	211	—	1.45	1.25
18	—	3056	363	—	1.46	1.26
20	—	7276	676	—	1.47	1.27
22	—	18094	1219	—	1.48	1.28
24	—	39135	1641	—	1.48	1.26

How to Find a Heuristic

General Strategy:

- Simplify the problem
- Compute the exact solution for the simplified problem
- Use the solution cost as heuristic

How to Find a Heuristic

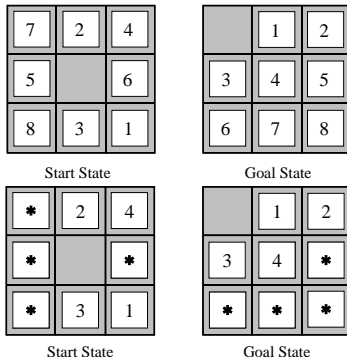
General Strategy:

- Simplify the problem
- Compute the exact solution for the simplified problem
- Use the solution cost as heuristic

For example:

- h_1 is the solution cost for the simplified 8-puzzle where tiles can be placed at an arbitrary position with a single action.
- h_2 corresponds to the exact solution, if tiles can be moved to an arbitrary position but actions are restricted to moving a tile to a neighboring position.

Pattern Databases



Idea: Compute the exact solution for each *pattern* with four numbers and use that value as heuristic. When more than one pattern applies, use the maximum value.
Better than Manhattan!

function IDA*(*problem*) **returns** a solution sequence

inputs: *problem*, a problem

static: *f-limit*, the current *f*- COST limit
root, a node

root \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

f-limit \leftarrow *f*- COST(*root*)

loop do

solution, *f-limit* \leftarrow DFS-CONTOUR(*root*, *f-limit*)

if *solution* is non-null **then return** *solution*

if *f-limit* = ∞ **then return** failure; **end**

function DFS-CONTOUR(*node*, *f-limit*) **returns** a solution sequence and a new *f*- COST limit

inputs: *node*, a node

f-limit, the current *f*- COST limit

static: *next-f*, the *f*- COST limit for the next contour, initially ∞

if *f*- COST[*node*] > *f-limit* **then return** null, *f*- COST[*node*]

if GOAL-TEST[*problem*](STATE[*node*]) **then return** *node*, *f-limit*

for each node *s* **in** SUCCESSORS(*node*) **do**



solution, *new-f* \leftarrow DFS-CONTOUR(*s*, *f-limit*)

if *solution* is non-null **then return** *solution*, *f-limit*

next-f \leftarrow MIN(*next-f*, *new-f*); **end**

return null, *next-f*

SMA* 1



function SMA*(*problem*) **returns** a solution sequence

inputs: *problem*, a problem

static: *Queue*, a queue of nodes ordered by *f*-cost

Queue \leftarrow MAKE-QUEUE({ MAKE-NODE(INITIAL-STATE[*problem*]) })

loop do

if *Queue* is empty **then return** failure

n \leftarrow deepest least-*f*-cost node in *Queue*

if GOAL-TEST(*n*) **then return** success

s \leftarrow NEXT-SUCCESSOR(*n*)

if *s* is not a goal and is at maximum depth **then**

f(*s*) $\leftarrow \infty$

else

f(*s*) \leftarrow MAX(*f*(*n*), *g*(*s*)+*h*(*s*))

if all of *n*'s successors have been generated **then**

 update *n*'s *f*-cost and those of its ancestors if necessary

if SUCCESSORS(*n*) all in memory **then** remove *n* from *Queue*

if memory is full **then**

 delete shallowest, highest-*f*-cost node in *Queue*


 remove it from its parent's successor list

 insert its parent on *Queue* if necessary

 insert *s* on *Queue*

end

Hill Climbing



function HILL-CLIMBING(*problem*) **returns** a solution state

inputs: *problem*, a problem

static: *current*, a node

next, a node

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

loop do

next \leftarrow a highest-valued successor of *current*

if VALUE[*next*] < VALUE[*current*] **then return** *current*

current \leftarrow *next*

end

Simulated Annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

inputs: *problem*, a problem

schedule, a mapping from time to “temperature”

static: *current*, a node

next, a node

T, a “temperature” controlling the probability of downward steps

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

for $t \leftarrow 1$ **to** ∞ **do**

T \leftarrow *schedule*[*t*]

if $T=0$ **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE[*next*] – VALUE[*current*]

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

Ridges

