

Problem Solving Using Search

Introduction to Artificial Intelligence

G. Lakemeyer

Winter Term 2018/19

Problem Solving Using Search

Here we consider **goal-oriented agents**.

Given an **initial (world) state**, the agent wants to reach a certain **goal** using appropriate **actions**, where actions transform one state into another.

A **goal** is a set of world states, which the agent finds desirable (wants to reach one of them).

Example: Driving from City A to City B.

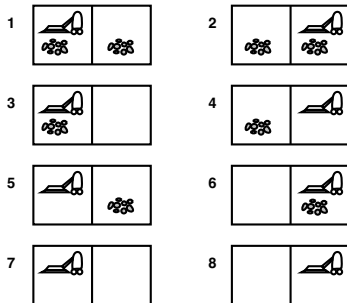
Problem formulation

- How to represent world states?
- Which actions should be considered?

Search: Finding an action sequence which transforms an initial state into a goal state.

Problem Classification

The vacuum-world states:



- **Single-state problem**
complete world knowledge,
complete knowledge about the actions
- **Multiple-state problem**
incomplete world knowledge,
complete knowledge about the actions
- **Contingency problem**
incomplete knowledge about actions,
needs to gather information at run-time
- **Exploration problem**
World states and effect of actions are
both unknown. Very difficult!

→ e.g. Mars Rover

Some Terminology (1)

Initial state: World state which the agent believes to be in initially.

State space: Set of all possible states.

Operator: Description of which state is reached by an action from a given state.

Successor function S : $S(x)$ returns the set of states reachable by any action from state x .

Goal test: Tests whether a state is a goal state.

Some Terminology (2)

Path: Sequence of actions.

Path cost: Cost function over paths, usually the sum of the costs of the actions of the path, often denoted as g .

Solution: Path from the initial state to a state that satisfies the goal test.

Search cost: Time and memory needed to find a solution.

Total cost: Sum of the path cost (of the solution) and search cost.

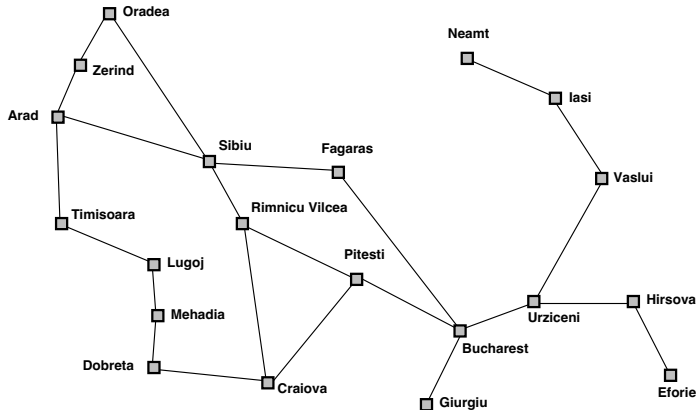
Note: For multiple-state problems, replace state by **state set**, operators apply to all elements of the set, and paths connect sets of states.

Choosing a State Space

It is an art to choose the right state space. Want to leave out unnecessary detail (**Abstraction**)!

Example: Drive from A to B.

Only represent the most important places.



8-Puzzle

only half the states are reachable,
i.e. $\frac{9!}{2}$ reach. states = 181 444

24-puzzle
(5x5)

has about
 10^{25} states

5	4	
6	1	8
7	3	2

Start State

1	2	3
8		4
7	6	5


Goal State

- **States:** Describe the location of each tile including the blank.
- **Operators:** Blank moves left, right, up, or down.
- **Goal test:** Does the state match picture on the right?
- **Path cost:** Each step costs 1 unit.
- **Search cost:** Problem is NP-complete in the size of the puzzle.

to find the shortest solution.

8-Queens Problem

- **Goal test:** 8 queens on the board, all safe.
- **Path cost:** 0 (only the solution counts).




Representation 1

- **States:** any placement of 0–8 queens.
- **Operators:** place a queen on the board.
- **Problem:** Too many states (64^8)!

If we allow at most 1 queen per square
then — $64 \times 63 \times \dots \times 57 \approx 3 \times 10^{14}$

8-Queens Problem

- **Goal test:** 8 queens on the board, all safe.
- **Path cost:** 0 (only the solution counts).



Representation 1


- **States:** any placement of 0–8 queens.
- **Operators:** place a queen on the board.
- **Problem:** Too many states (64^8)!

Representation 2

- **States:** 0–8 on board, all safe.
- **Operators:** place a queen starting from left.
- **Problem:** few state (2057), but sometimes deadend (see figure).

8-Queens Problem

- **Goal test:** 8 queens on the board, all safe.
- **Path cost:** 0 (only the solution counts).



Representation 1

- **States:** any placement of 0–8 queens.
- **Operators:** place a queen on the board.
- **Problem:** Too many states (64^8)!

Representation 2

- **States:** 0–8 on board, all safe.
- **Operators:** place a queen starting from left.
- **Problem:** few state (2057), but sometimes deadend (see figure).


Representation 3

- **States:** 8 queens on board, one per column.
- **Operators:** move a queen under attack in the same column.

8^8 state

Vacuum World 1


Complete World Knowledge (single-state problem):



- **States:** All combinations of robot location and presence of dust (8).
- **Operators:** left (L), right (R), suck (S).
- **Goal test:** no dust in any room.
- **Path cost:** 1 unit per action.

Vacuum World 2


Incomplete Knowledge (multiple-state problem).



- **State set:** All subsets of the 8 states.
- **Goal test:** No dust anywhere in the elements of the state set.

Missionaries and cannibals (MaC)

John McCarthy



3 missionaries and 3 cannibals want to cross a river. There is a row boat holding at most 2 people. Should the cannibals ever outnumber the missionaries, the missionaries will be eaten. How do all six make it across safely.

Non-Intended MaC Solutions

- ❶ All six walk across the bridge nearby.
- ❷ They take a helicopter.
- ❸ There are three more missionaries on the other side.
- ❹ ...

Contradicts the underlying assumption that the description of the puzzle contains **all relevant** information.

Gricean principle

Sabotaging the Correct MaC Solution

- 1 There is a hole in the boat.
- 2 The missionaries are too weak to row.
- 3 One missionary goes overboard and drowns.
- 4 ...

Contradicts the assumptions of what happens **normally**.

Formalizing the MaC-Problem

- **States:** Triple (x,y,z) with $0 \leq x, y \leq 3$ and $0 \leq z \leq 1$, where x , y , and z denote how many missionaries, cannibals, and boats are currently on the left river bank.
- **Initial state:** $(3,3,1) \xrightarrow{2 \text{ cannibals}} (3,1,0)$
- **Operators:** Either one missionary, one cannibal, two missionaries, two cannibals, or one of each cross the river in a boat. (i.e. 5 operators)
- **Note:** not every state is reachable [e.g. $(0,0,1)$].
- **Goal state:** $(0,0,0)$
- **Path cost:** 1 unit per river crossing.

Note: Solution due to Saul Amarel (1968).

Search in General

Starting from the initial state we always **expand** one state, i.e. we create all successor states of that state.


Search tree = the tree induced by expansion.

Nodes contain world states;
edges = actions.


(a) The initial state

Arad

(b) After expanding Arad

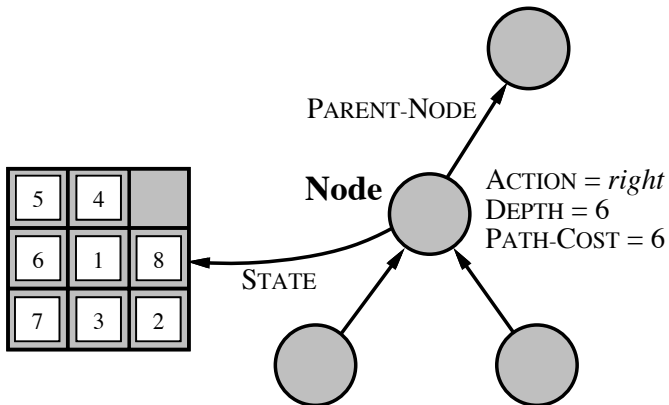


(c) After expanding Sibiu



```
function GENERAL-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

Search Nodes



Evaluating Search Strategies

- **Completeness:** Does it always find a solution if one exists?
- **Time Complexity:** How long does it take to find a solution in the worst case?
- **Space Complexity:** How much memory is used in the worst case?
- **Optimality:** Does it always find the best solution?

Uninformed vs. Informed Search

Uninformed (blind) search:

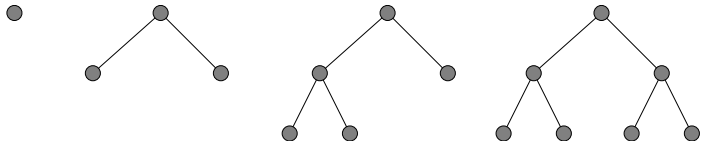
- No information available about the length or the cost of a solution.
- The main methods are:
breadth-first search, uniform cost search, depth-first search,
depth-limited search, iterative deepening, bidirectional search.

Informed (heuristic) search:

- Here the user has information about the length or the cost of a solution, which often helps in speeding up the search.
- We will look at:
Greedy search, A^* and some variants, hill-climbing, and simulated annealing.

Breadth-First Search

Expand the nodes in the same order in which they are generated.



- Always finds the shallowest solution (completeness).
- Solutions are optimal, provided all actions have identical non-negative cost.

Cost of Breadth-First Search

time = space complexity!

The Cost is very high!. Let b be the max. branching factor, d the depth of a solution. Then at most

$$b + b^2 + b^3 + \dots + b^d + b^{d+1} - b$$

nodes need to be generated, i.e. $O(b^{d+1})$.


→ depends on doing goal test only when choosing a node for expansion

Example: $b = 10$, 10000 nodes/s; 1000bytes/node:

Depth	Nodes	Time	Memory
2	1100	.11 seconds	1 megabyte
4	111,100	11 seconds	106 megabytes
6	10^7	19 minutes	10 gigabytes
8	10^9	31 hours	1 terabytes
10	10^{11}	129 days	101 terabytes
12	10^{13}	35 years	10 petabytes
14	10^{15}	3,523 years	1 exabyte

Uniform cost search

A modified breadth-first search, which always expands node n with the least path cost $g(n)$.



Note: Always finds the cheapest solution, provided

$$g(\text{successor}(n)) \geq g(n) \text{ for all } n.$$

optimality guaranteed for non-negative action costs.

Depth-First Search


not complete if depth is unbounded.
not optimal

Always expand the node at the deepest level.

Space: $O(b \times m)$

m : maximal depth

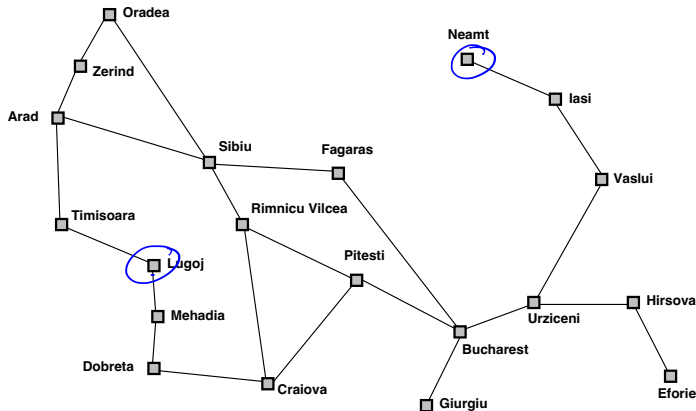
time: $O(b^m)$



Depth-limited Search

Perform depth-first search, but only up to a pre-specified depth.

Route planning: For n cities the maximal depth is $n - 1$.



Actually, a maximal depth of 9 suffices (called the **diameter** of the problem).

(difficult to find in general)

Iterative Deepening 1

- Combines depth- and breadth-first search.
- Optimal and complete like breadth-first search, but requires less space.


```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution sequence  
  inputs: problem, a problem  
  
  for depth  $\leftarrow$  0 to  $\infty$  do  
    if DEPTH-LIMITED-SEARCH(problem, depth) succeeds then return its result  
  end  
  return failure
```

Limit = 0 ●

Limit = 1 ●

Limit = 2 ●

Limit = 3 ●



Space $O(b \times d)$
Time $O(b^d)$?

Iterative Deepening 2

Number of expansions:

$$(d)b + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

*children of the root generated d times
nodes at level d*

Compared to breadth-first search: $b + b^2 + b^3 + \dots + b^{d-1} + b^d + b^{d+1} - b$.

Note: iterative deepening is better!

Example: $b = 10$, $d = 5$

Breadth-first search:

$$10 + 100 + 1000 + 10000 + 100000 + 999990 = 1111110$$

Iterative deepening:

$$50 + 400 + 3000 + 20000 + 100000 = 123450$$


*if BFS does a goal test
when a node is generated, then
BFS is a little better (111110)*

Time complexity: $O(b^d)$

Space complexity: $O(b \times d)$

Note: Preferred method if search depth unknown!

Bidirectional Search



Assuming that forward and backward search are symmetric, we can ideally obtain search times of $O(2 \times b^{d/2}) = O(b^{d/2})$.

Example: For $b = 10$, $d = 6$ we have 2222 nodes instead of 1111111!

Problems with Bidirectional Search

- Operators not always reversible. (Need to compute the predecessor node.)
- Sometimes there are very many goal states, which are described only incompletely. E.g., what are the predecessors of “checkmate” in chess?
- We need efficient methods to decide when the two search methods have met.
- Which kind of search method should one use for each direction? (In the figure it is breadth-first search for both, but that is not always optimal.)

Comparison

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^{d+1})$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^{d+1})$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

b = maximal branching factor

d = depth of a solution path

m = maximal search depth

ℓ = depth restriction

assumes action cost $\geq \epsilon > 0$
 $C^* \triangleq$ optimal cost to a goal.

^a: complete if b is finite; ^b: complete if step costs $\geq \epsilon > 0$;

^c: optimal if step costs identical; ^d: if both directions use BFS.