

# Machine Learning - Exercise 4

Companion Slides (adapted from Lucas Beyer)

Paul Voigtlaender   Sabarinath Mahadevan

December 6, 2018

# What's the plan?

- ▶ Exercise overview
- ▶ Deep learning in a nutshell
- ▶ Backprop in detail

# Overview

- ▶ Goal: implement a simple DL framework from scratch
- ▶ Tasks:
  - ▶ Compute derivatives (Jacobians)
  - ▶ Write code

# Deep Learning in a Nutshell

Given:

- ▶ Training data  $X = \{x_i\}_{i=1\dots N}$  with  $x_i \in \mathbb{I}$ , usually as  $X \in \mathbb{R}^{N \times N_I}$
- ▶ Training labels  $T = \{t_i\}_{i=1\dots N}$  with  $t_i \in \mathbb{O}$

Choose

- ▶ Parameterized, (sub-)differentiable function  $F(X, \theta) : \mathbb{I} \times \mathbb{P} \rightarrow \mathbb{O}$ , with
  - ▶ typically, input-space  $\mathbb{I} = \mathbb{R}^{N_I}$  (generic data),  $\mathbb{I} = \mathbb{R}^{3 \times H \times W}$  (images), ...
  - ▶ typically, output-space  $\mathbb{O} = \mathbb{R}^{N_O}$  (regression),  $\mathbb{O} = [0, 1]^{N_O}$  (probabilistic classification), ...
  - ▶ typically, parameter-space  $\mathbb{P} = \mathbb{R}^{N_P}$
- ▶ (Sub-)differentiable criterion/loss  $\mathcal{L}(T, F(X, \theta)) : \mathbb{O} \times \mathbb{O} \rightarrow \mathbb{R}$

RGB\*Height\*Width

Find:

$$\theta^* = \underset{\theta \in \mathbb{P}}{\operatorname{argmin}} \mathcal{L}(T, F(X, \theta))$$

Assumption:

$$\mathcal{L}(T, F(X, \theta)) = \frac{1}{N} \sum_{i=1}^N \ell(t_i, F(x_i, \theta))$$



Visual Computing  
Institute

RWTH AACHEN  
UNIVERSITY

# Deep Learning in a Nutshell

Given:

- ▶ Training data  $X = \{x_i\}_{i=1\dots N}$  with  $x_i \in \mathbb{I}$ , usually as  $X \in \mathbb{R}^{N \times N_I}$
- ▶ Training labels  $T = \{t_i\}_{i=1\dots N}$  with  $t_i \in \mathbb{O}$

Choose

- ▶ Parameterized, (sub-)differentiable function  $F(X, \theta) : \mathbb{I} \times \mathbb{P} \rightarrow \mathbb{O}$ , with
  - ▶ typically, input-space  $\mathbb{I} = \mathbb{R}^{N_I}$  (generic data),  $\mathbb{I} = \mathbb{R}^{3 \times H \times W}$  (images), ...
  - ▶ typically, output-space  $\mathbb{O} = \mathbb{R}^{N_O}$  (regression),  $\mathbb{O} = [0, 1]^{N_O}$  (probabilistic classification), ...
  - ▶ typically, parameter-space  $\mathbb{P} = \mathbb{R}^{N_P}$
- ▶ (Sub-)differentiable criterion/loss  $\mathcal{L}(T, F(X, \theta)) : \mathbb{O} \times \mathbb{O} \rightarrow \mathbb{R}$

RGB\*Height\*Width

Find:

$$\theta^* = \underset{\theta \in \mathbb{P}}{\operatorname{argmin}} \mathcal{L}(T, F(X, \theta))$$

Assumption:

$$\mathcal{L}(T, F(X, \theta)) = \frac{1}{N} \sum_{i=1}^N \ell(t_i, F(x_i, \theta))$$



Visual Computing  
Institute

**RWTH**AACHEN  
UNIVERSITY

# Deep Learning in a Nutshell

Given:

- ▶ Training data  $X = \{x_i\}_{i=1\dots N}$  with  $x_i \in \mathbb{I}$ , usually as  $X \in \mathbb{R}^{N \times N_I}$
- ▶ Training labels  $T = \{t_i\}_{i=1\dots N}$  with  $t_i \in \mathbb{O}$

Choose

- ▶ Parameterized, (sub-)differentiable function  $F(X, \theta) : \mathbb{I} \times \mathbb{P} \rightarrow \mathbb{O}$ , with
  - ▶ typically, input-space  $\mathbb{I} = \mathbb{R}^{N_I}$  (generic data),  $\mathbb{I} = \mathbb{R}^{3 \times H \times W}$  (images), ...
  - ▶ typically, output-space  $\mathbb{O} = \mathbb{R}^{N_O}$  (regression),  $\mathbb{O} = [0, 1]^{N_O}$  (probabilistic classification), ...
  - ▶ typically, parameter-space  $\mathbb{P} = \mathbb{R}^{N_P}$
- ▶ (Sub-)differentiable criterion/loss  $\mathcal{L}(T, F(X, \theta)) : \mathbb{O} \times \mathbb{O} \rightarrow \mathbb{R}$

RGB\*Height\*Width

Find:

$$\theta^* = \underset{\theta \in \mathbb{P}}{\operatorname{argmin}} \mathcal{L}(T, F(X, \theta))$$

Assumption:

$$\mathcal{L}(T, F(X, \theta)) = \frac{1}{N} \sum_{i=1}^N \ell(t_i, F(x_i, \theta))$$



Visual Computing  
Institute

**RWTH**AACHEN  
UNIVERSITY

# Deep Learning in a Nutshell

Given:

- ▶ Training data  $X = \{x_i\}_{i=1\dots N}$  with  $x_i \in \mathbb{I}$ , usually as  $X \in \mathbb{R}^{N \times N_I}$
- ▶ Training labels  $T = \{t_i\}_{i=1\dots N}$  with  $t_i \in \mathbb{O}$

Choose

- ▶ Parameterized, (sub-)differentiable function  $F(X, \theta) : \mathbb{I} \times \mathbb{P} \rightarrow \mathbb{O}$ , with
  - ▶ typically, input-space  $\mathbb{I} = \mathbb{R}^{N_I}$  (generic data),  $\mathbb{I} = \mathbb{R}^{3 \times H \times W}$  (images), ...
  - ▶ typically, output-space  $\mathbb{O} = \mathbb{R}^{N_O}$  (regression),  $\mathbb{O} = [0, 1]^{N_O}$  (probabilistic classification), ...
  - ▶ typically, parameter-space  $\mathbb{P} = \mathbb{R}^{N_P}$
- ▶ (Sub-)differentiable criterion/loss  $\mathcal{L}(T, F(X, \theta)) : \mathbb{O} \times \mathbb{O} \rightarrow \mathbb{R}$

RGB\*Height\*Width

Find:

$$\theta^* = \underset{\theta \in \mathbb{P}}{\operatorname{argmin}} \mathcal{L}(T, F(X, \theta))$$

Assumption:

$$\mathcal{L}(T, F(X, \theta)) = \frac{1}{N} \sum_{i=1}^N \ell(t_i, F(x_i, \theta))$$



Visual Computing  
Institute

RWTH AACHEN  
UNIVERSITY

# Backprop

Take derivative with respect to  $\theta$

$$\begin{aligned} D_{\theta} \frac{1}{N} \sum_{i=1}^N \ell(t_i, F(x_i, \theta)) &= \frac{1}{N} D_{\theta} \sum_{i=1}^N \ell(t_i, F(x_i, \theta)) \\ &= \frac{1}{N} \sum_{i=1}^N D_{\theta} F(x_i, \theta) \cdot D_F \ell(t_i, F(x_i, \theta)) \end{aligned}$$

Chain rule

Assumption:

$F$  is hierarchical:  $F(x_i, \theta) = f_1(f_2(f_3(\dots x_i \dots, \theta_3), \theta_2), \theta_1)$

$D_{\theta_1} F(x_i, \theta) = D_{\theta_1} f_1(f_2, \theta_1)$

$D_{\theta_2} F(x_i, \theta) = D_{\theta_2} f_2(f_3, \theta_2) \cdot D_{f_2} f_1(f_2, \theta_1)$

$D_{\theta_3} F(x_i, \theta) = D_{\theta_3} f_3(\dots, \theta_3) \cdot D_{f_3} f_2(f_3, \theta_2) \cdot D_{f_2} f_1(f_2, \theta_1)$

Where  $f_2 = f_2(f_3(\dots x_i \dots, \theta_3), \theta_2)$  etc.



# Backprop

Take derivative with respect to  $\theta$

$$\begin{aligned} D_{\theta} \frac{1}{N} \sum_{i=1}^N \ell(t_i, F(x_i, \theta)) &= \frac{1}{N} D_{\theta} \sum_{i=1}^N \ell(t_i, F(x_i, \theta)) \\ &= \frac{1}{N} \sum_{i=1}^N D_{\theta} F(x_i, \theta) \cdot D_F \ell(t_i, F(x_i, \theta)) \end{aligned}$$

Chain rule

Assumption:

$F$  is hierarchical:  $F(x_i, \theta) = f_1(f_2(f_3(\dots x_i \dots, \theta_3), \theta_2), \theta_1)$

$D_{\theta_1} F(x_i, \theta) = D_{\theta_1} f_1(f_2, \theta_1)$

$D_{\theta_2} F(x_i, \theta) = D_{\theta_2} f_2(f_3, \theta_2) \cdot D_{f_2} f_1(f_2, \theta_1)$

$D_{\theta_3} F(x_i, \theta) = D_{\theta_3} f_3(\dots, \theta_3) \cdot D_{f_3} f_2(f_3, \theta_2) \cdot D_{f_2} f_1(f_2, \theta_1)$

Where  $f_2 = f_2(f_3(\dots x_i \dots, \theta_3), \theta_2)$  etc.

# Jacobians

The loss:

$$D_F \ell(t_i, F(x_i, \theta)) = \begin{pmatrix} \partial_{F_1} \ell \\ \vdots \\ \partial_{F_{N_F}} \ell \end{pmatrix} \in \mathbb{R}^{N_F \times 1}$$

The functions (modules):

$$f(z, \theta) = (f_1(z_1, \dots, z_{N_z}; \theta) \quad \dots \quad f_{N_f}(z_1, \dots, z_{N_z}; \theta)) \in \mathbb{R}^{1 \times N_f}$$

$$D_z f(z, \theta) = \begin{pmatrix} \partial_{z_1} f_1 & \dots & \partial_{z_1} f_{N_f} \\ \vdots & \ddots & \vdots \\ \partial_{z_{N_z}} f_1 & \dots & \partial_{z_{N_z}} f_{N_f} \end{pmatrix} \in \mathbb{R}^{N_z \times N_f}$$

The loss:

$$D_F \ell(t_i, F(x_i, \theta)) = \begin{pmatrix} \partial_{F_1} \ell \\ \vdots \\ \partial_{F_{N_F}} \ell \end{pmatrix} \in \mathbb{R}^{N_F \times 1}$$

The functions (modules):

$$f(z, \theta) = (f_1(z_1, \dots, z_{N_z}; \theta) \quad \dots \quad f_{N_f}(z_1, \dots, z_{N_z}; \theta)) \in \mathbb{R}^{1 \times N_f}$$

$$D_z f(z, \theta) = \begin{pmatrix} \partial_{z_1} f_1 & \dots & \partial_{z_1} f_{N_f} \\ \vdots & \ddots & \vdots \\ \partial_{z_{N_z}} f_1 & \dots & \partial_{z_{N_z}} f_{N_f} \end{pmatrix} \in \mathbb{R}^{N_z \times N_f}$$

# Modules

Looking at module  $f_2$ :

$$D_{\theta_3} F(x_i, \theta) = [D_{\theta_3} f_3(\dots, \theta_3)] \underbrace{[D_{f_3} f_2(\overbrace{f_3}^{\text{input}}, \theta_2)]}_{\text{Jacobian wrt. input}} \underbrace{[D_{f_2} f_1(f_2, \theta_1)]}_{\text{grad\_output}}$$

grad\_input

Three (core) functions per module:

fprop: compute the output  $f_i(z, \theta_i)$  given the input  $z$  and current parameters  $\theta_i$

grad\_input: compute  $D_z f_i(z, \theta_i) \cdot \text{grad\_output}$

grad\_param: compute  $\nabla_{\theta_i} = D_{\theta_i} f_i(z, \theta_i) \cdot \text{grad\_output}^T$

Typically:

fprop caches its input and/or output for later reuse

grad\_input and grad\_param are combined into single bprop function to share computation

# Usage/Training

```
1:  $net = [f1, f2, \dots, f_{N_f}], \ell = \text{criterion}$ 
2: for  $Xb, Tb$  in batched  $X, T$  do
3:    $z = Xb$ 
4:   for module in net do
5:      $z = \text{module.fprop}(z)$ 
6:   end for
7:    $\text{costs} = \ell.\text{fprop}(z, Tb)$ 
8:    $\partial z = \ell.\text{bprop}([\frac{1}{N_B} \dots \frac{1}{N_B}])$ 
9:   for module in reversed(net) do
10:     $\partial z = \text{module.bprop}(\partial z)$ 
11:   end for
12:   for module in net do
13:     $\theta, \partial \theta = \text{module.params}(), \text{module.grads}()$ 
14:     $\theta = \theta - \lambda \cdot \partial \theta$ 
15:   end for
16: end for
```

## Example: Linear aka. Fully-connected module

$$f(z, W, b) = z \cdot W + b \in \mathbb{R}^{1 \times N_f}$$

Where  $z \in \mathbb{R}^{1 \times N_z}$ ,  $W \in \mathbb{R}^{N_z \times N_f}$ ,  $b \in \mathbb{R}^{1 \times N_f}$ , and  
 $\text{grad\_output} = D_f \ell(f(z, W, b))^T \in \mathbb{R}^{N_f \times 1}$

The gradients are

- ▶  $\mathbb{R}^{N_z \times N_f} \ni \text{grad\_W} = z^T \cdot \text{grad\_output}^T$
- ▶  $\mathbb{R}^{N_f \times 1} \ni \text{grad\_b} = \text{grad\_output}$
- ▶  $\mathbb{R}^{N_z \times 1} \ni \text{grad\_input} = W \cdot \text{grad\_output}$

# Gradient Checking

Crucial debugging method!

Compare Jacobian computed by finite differences using the fprop function to Jacobian computed by the bprop function.

Advice: Use (small) random input  $x$ , and  $h_i = \sqrt{\epsilon ps} \max(x_i, 1)$ .

Finite-difference: first column of Jacobian as:

$$x_- = (x_1 - h_1 \quad x_2 \quad x_3 \quad \dots \quad x_{N_x})$$

$$x_+ = (x_1 + h_1 \quad x_2 \quad x_3 \quad \dots \quad x_{N_x})$$

$$J_{\bullet,1} = \frac{\text{fprop}(x_+) - \text{fprop}(x_-)}{2h_1}$$

Backprop: first row of Jacobian as:

$$\text{fprop}(x)$$

$$J_{1,\bullet} = \text{bprop}(1 \quad 0 \quad 0 \quad \dots \quad 0)$$



# Mini-Batching

Linear layer (without mini-batching)

$$f(z, W, b) = z \cdot W + b$$

$$z \in \mathbb{R}^{1 \times N_z}, W \in \mathbb{R}^{N_z \times N_f}, b \in \mathbb{R}^{1 \times N_f}$$

Stack  $z$  into mini-batch matrix with batch size  $N \Rightarrow z \in \mathbb{R}^{N \times N_z}$

Now the multiplication  $z \cdot W$  can be performed for all examples in one pass and  $b$  can be added by broadcasting (repeating)  $b$  to  $\hat{b}$

$$f(z, W, b) = \begin{pmatrix} \text{---} & f_1 & \text{---} \\ \vdots & \vdots & \vdots \\ \text{---} & f_N & \text{---} \end{pmatrix} = \begin{pmatrix} z_{1,1} & \cdots & z_{1,N_z} \\ \vdots & \ddots & \vdots \\ z_{N,1} & \cdots & z_{N,N_z} \end{pmatrix} \cdot \begin{pmatrix} W_{1,1} & \cdots & W_{1,N_f} \\ \vdots & \ddots & \vdots \\ W_{N_z,1} & \cdots & W_{N_z,N_f} \end{pmatrix} + \begin{pmatrix} \text{---} & b & \text{---} \\ \vdots & \vdots & \vdots \\ \text{---} & b & \text{---} \end{pmatrix}$$



# Rule-of-thumb result on MNIST

Linear( $28 \times 28, 10$ ), Softmax should give  $\pm 750$  errors.

Linear( $28 \times 28, 200$ ), tanh, Linear( $200, 10$ ), SoftMax should give  $\pm 250$  errors.

Typical learning rates  $\lambda \in [0.1, 0.01]$

Typical batch-sizes  $N_B \in [100, 1000]$

Initialize weights as  $\mathbb{R}^{M \times N} \ni W \sim \mathcal{N}(0, \sigma = \sqrt{\frac{2}{M+N}})$  and  $b = 0$

Don't forget data pre-processing, here at least divide values by 255 (max pixel value).